



A Deep Learning Approach to American Sign Language Character Classification

By Arvind Shankar, Ryan Leavitt,
James Dika, Stephen Young

Data Acquisition and Preprocessing

- ◆ Data Acquisition:
 - ◆ Class data – 1844 images
 - ◆ MNIST data – added 13,013 more images
- ◆ Preprocessing:
 - ◆ Integer label encoding
 - ◆ Gray scaling provided data
 - ◆ Reduces dimensionality by factor of 3
 - ◆ Standardizing grayscale data
 - ◆ This data has mean 0, unit variance
 - ◆ Training and Test Split
- ◆ Other preprocessing methods performed poorly in validation
 - ◆ PCA, HOG, and LDA

Model Selection

- ◆ In Lab 3 we trained and validated these classifiers:
 - ◆ LDA, KNN, Decision Trees, Random Forests, SVM, and MLPs
- ◆ The highest accuracy classifier was an MLP with 74% test accuracy
- ◆ Then we investigated deep learning options for the classifier and decided to go with a Convolutional Neural Net for our model

	Algo	Accuracies
0	LDA	0.308943
1	KNN	0.634146
2	Decision Tree	0.346883
3	Random Forest	0.430894
4	SVM	0.092141
5	MLP	0.739837

Our CNN

- ◆ Following online tutorials and other resources, we created a CNN class with the following parameters kept constant:
 - ◆ Output activation function
 - ◆ Softmax activation function since this is a multiclass problem
 - ◆ Softmax activation outputs a distribution of probabilities – the index with the highest probability is the predicted label
 - ◆ Objective function
 - ◆ Literature also revealed that minimizing the Cross-Entropy loss function works well along with the Softmax function
 - ◆ Learning Rate
 - ◆ Initialized to 0.001 based on PyTorch documentation
 - ◆ Gradient Descent Optimizer
 - ◆ Used Adam since it is the most common method used
 - ◆ This results in an adaptive learning rate

$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

$$\text{loss}(x, \text{class}) = -\log \left(\frac{\exp(x[\text{class}])}{\sum_j \exp(x[j])} \right) = -x[\text{class}] + \log \left(\sum_j \exp(x[j]) \right)$$

<https://pytorch.org/docs/stable/nn.html>

Tuning CNN Parameters

- ◆ We did testing to find optimal values for the following parameters:
 - ◆ Number of convolutional layers
 - ◆ Kernel size
 - ◆ Number of inputs to first linear layer
 - ◆ Training set
 - ◆ Batch size
 - ◆ Number of epochs
- ◆ Testing methodology:
 - ◆ Train CNNs with different values for the parameter in question, other parameters kept constant
 - ◆ Score CNNs on validation partition of the grayscale provided dataset – representative of the real test set
 - ◆ Keep parameter values which maximize validation accuracy

Tuning The CNN cont.

- ◆ Kernel Size

- ◆ Tested kernel sizes 3 and 5, found that 3x3 was better

- ◆ Convolutional Layers

- ◆ Tested 3 and 4 convolutional layers – 3 layers performed much better

- ◆ Number of Inputs to First Linear Layer

- ◆ Reducing number of outputs in last convolutional layer lead to much worse performance

- ◆ Training Data

- ◆ 3 datasets – grayscale provided set, MNIST combined with grayscale set, grayscale & standardized provided set
 - ◆ MNIST data didn't help – makes sense since the MNIST set is not representative of the real test data
 - ◆ Standardized dataset performed very poorly
 - ◆ The model performed best with data with just the grayscale provided dataset

Tuning The CNN cont.

◆ Batch Size and Epochs


- ◆ Tested 19 different combinations with small to large values for both parameters – found that batch size 128 and 1000 epochs gave us the best model
- ◆ More epochs might lead to increased accuracy, but execution time was already very long – this made it difficult to do more tests

◆ Examining training and validation accuracy over 1000 epochs:



Results

- ◇ The best model we obtained had an overall accuracy of 89.7%
- ◇ The confusion matrix and the accuracy for each letter for this model are shown below:



Letter	Accuracy	Letter	Accuracy
A	89.1%	F	92.7%
B	90.2%	G	81.8%
C	92.9%	H	92.1%
D	90.0%	I	88.2%
E	90.9%		

Possible Improvements

- ◆ There were a few parameters we didn't vary that could potentially be optimized through further testing:
 - ◆ Loss function
 - ◆ There are others besides cross-entropy that work with multiclass problems
 - ◆ Optimizer
 - ◆ AdaGrad and RMSProp are alternative adaptive learning rate optimizers
- ◆ Could also have researched alternative preprocessing methods for feature extraction generation besides the ones we tried in Lab 2

Conclusions

- ◆ We made a CNN with 3 convolutional layers using 3x3 kernels, 2D pooling layers, and 2 linear layers
- ◆ We determined most parameters for our model through incremental experimentation while the constant parameter choices were based on the literature and resources
- ◆ Training this CNN on the grayscale provided data with a batch size of 128 for 1000 epochs, we used the Adam optimizer to minimize cross-entropy loss
- ◆ Testing this CNN on a realistic test set, we got an overall accuracy of 89.7%, with 6 out of the 9 letters being classified with at least 90% accuracy